



SAPIENZA
UNIVERSITÀ DI ROMA

DEPARTMENT OF INFORMATION ENGINEERING, ELECTRONICS AND
TELECOMMUNICATIONS (DIET)

Challenge #1 - Topology Design


NBD - NETWORKING FOR BIG DATA AND LABORATORY

Professor: Andrea Baiocchi	Group Viterbi:	
	Giovanni Giunta	1177327
	Mehrzad Jafari Ranjbar	1937944
	Borbála Sára Tóth Gy	2066306

1 Assignment - Part 1


(i) We implemented three different Python functions to check the connectivity of Erdős–Rényi and r -regular random graphs, and then tested their complexity for different number of nodes K .

In the plots below, we can see that the Irreducibility and the Laplacian approach have the complexity of $O(K^3)$. However, the Irreducibility method is less efficient, as it requires performing several matrix multiplications, which is a computationally expensive task. The BFS algorithm is the most efficient among the three proposed methods, with a time complexity of $O(V + E)$. These results are consistent for both p-ER and r -regular random graphs.



(ii) In the bottom left plot, we observe the behavior of a p-ER random graph with $K = 100$ nodes at different levels of edge probability p . It appears that Erdős–Rényi random graphs are almost certainly connected even with a relatively low edge probability of 0.09.

(iii) In the bottom right plot we observe the probability for a R -regular random graph to be connected, as we increase the number of nodes K , in the case of $r = 2$ and $r = 8$. When $r = 8$, the graph is always connected, as the probability of connectivity is 1 for all values of K . When $r = 2$, the probability of connectivity fluctuates and is lower than when $r = 8$. As we can see, a lower degree of regularity can lead to a lower probability of connectivity as we increase the value of K due to fewer edges connecting the vertices.



2 Assignment - Part 2

(i) To evaluate the **Mean Response Time**, we ran a loop ranging from 1 to 10,000 servers. At each iteration N , we:

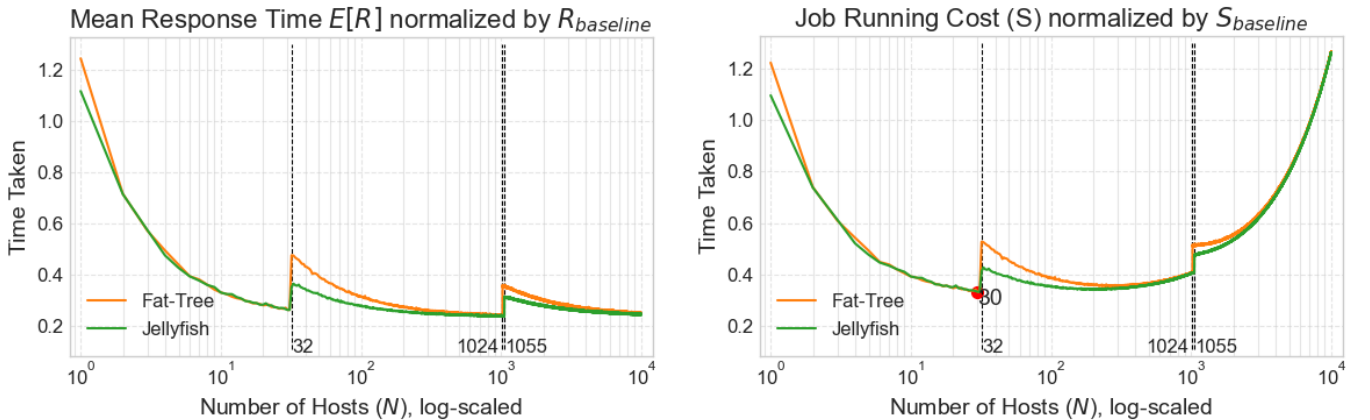
1. Create a dictionary in which the keys are the N closest servers to A , and the values are the number of hops from A to the server
2. Compute the average throughput, that is $TH_i = C \frac{1/T_i}{\sum_{j=1}^N 1/T_j}$ for each server i
3. Measure the time needed to send a fraction of data to each of the N servers: $\frac{(L_f + (L_f * f))/N}{TH_i}$
4. Run a simulation with 500 steps, where at each iteration we simulate:
 - (a) The time needed by each server i to compute its share of processing, equal to $T_0 + E[X_i]$. In short, we sum the given constant T_0 to samples of size N drawn from a negative exponential distribution with mean equal to $E[X_i] = E[X]/N$. Each result is appended to the **job computations** list after each iteration
 - (b) The fraction of data $L_{o,i}$ returned by each server i at the end of its share of job computation. We draw samples of size N from a uniform distribution defined in $[0, 2L_o/N]$. Each result is appended to the **outputs** list after each iteration
5. Compute the average of all the **job computations** samples (returns a single list)
6. Compute the average of all the **outputs** samples (returns a single list)
7. Add the overhead to each value of the **outputs** list, and measure the time needed to send each fraction of output back to server A : $\frac{(L_{o,i} + (L_{o,i} * f))/N}{TH_i}$
8. Sum the time needed to send the fraction of data from A to i (see 3.), the job computation time of i (see 5.), and the time needed to send the output from i back to A (see 7.). Now we have a single list with the average response times for each server i
9. Return the **Mean Response Time**, that is, the maximum value in the list with all the average response times (see 8.), since all the servers work in parallel

Selecting the N closest servers to A (and their distance from A in hops) is a trivial task, and can be easily computed. However, we proceeded nonetheless with generating a Fat-Tree and a Jellyfish network and performed a simulation starting from the actual graphs. In any case, both approaches lead to the same result. It is also worth mentioning that both the Fat-Tree and the Jellyfish networks have been given the same number of servers, so to compare them in a meaningful way.

In fact, given a number of ports $n = 64$, the Jellyfish network was built with **servers** = $\frac{n^3}{4}$ and **switches** = $\frac{\text{servers}}{r}$, with $r = \frac{n}{2}$. This way, we end up having 65,536 servers in both networks.

(ii) In the bottom left plot we have the results for the **Mean Response Time**. At first, it decreases dramatically as we distribute the job over the first 31 servers. Then it increases all of a sudden for $N = 32$, exactly when we start taking into account servers that are more distant from A . The effect is mitigated by adding more and more servers. For Fat-Tree we have an additional jump at $N = 1024$, when each additional server starts having $\text{hops} = 6$. For Jellyfish the jump is at $N = 1055$, when additional servers start having $\text{hops} = 4$. By adding more and more servers the mean response time decreases again, until it becomes almost flat towards $N = 10,000$: the job computation time becomes so small that T_0 becomes more relevant than $E[X_i]$.

(iii) In the bottom right plot we have the **Job Running Cost**. We can use it as a measure to explain the overall job computation cost of our networks, and is given by $S = E[R] + \xi E[\Theta]$. In this case, $E[R]$ and $E[\Theta]$ are in contrast with each other. In fact, we can observe that after each jump we improve our performance, but as we add more servers, $E[\Theta]$ prevails and the function increases. This happens because $E[\Theta]$ is computed as the sum of the computation times of all the N servers.



(iv) It seems that the optimal number of servers is the same for both the Fat-Tree and the Jellyfish networks, and it is equal to 30. It corresponds to the point of minimum of the **Job Running Cost** function, and it has been highlighted on the chart. It is worth noting, however, that slightly different results can be achieved with a different number of Monte Carlo simulations.

(v) In general, we can say that the Jellyfish network topology performs better than the Fat-Tree, both in terms of **Mean Response Time** and **Job Running Cost**. In both scenarios, it is evident that adding more and more servers is not always convenient. Distributing the job over many servers can be an advantage only if the distance from A (and consequently the throughput), is not significant. Adding to our pool of servers even one single node that is too distant from A , can dramatically decrease the overall performance of our network. Once the threshold is passed, bringing back the performance to acceptable levels requires too many additional resources, which leads to a waste of computational power, as we can see from the **Job Running Cost** plot: we would perhaps get close to the same response time, but at a higher job running cost.